

Live Session Agenda

Clever · Senior Product Manager take-home Joshua Agbabiaka Suggested length: 60 minutes

What I want out of this session

Three things, in order of importance:

1. **Correct my diagnosis.** I've read two CSV files. You've lived this. I expect at least one of my conclusions to be wrong, and I'd rather find out here than in week 3.
2. **Confirm the one number the whole business can share.** If Sales, Ops and Engineering leave with the same definition of a missed lead, most of the finger-pointing stops.
3. **Pressure-test the sequence.** Especially what I chose *not* to fix.

How I'd run the hour

Time	Segment	Purpose
0:00–0:05	Framing	The one-sentence diagnosis, then straight into questions
0:05–0:20	Questions I can't answer from the data	The seven blockers below
0:20–0:35	Walk the prototype	Show the ledger, take challenges live
0:35–0:50	The plan, and what I'd leave alone	Where I most expect disagreement
0:50–1:00	What week 1 looks like if I started Monday	Who I'd need, and for how long

I'd rather spend the first third asking than presenting. If the group wants to go straight to the prototype, I'll adapt.

1. Questions that change the answer

These seven affect conclusions in the proposal. I've ordered them by how much rides on them.

1. What timezone are received_at and attempted_at recorded in? Why it matters: I found 73 attempts outside 8am–9pm in the lead’s local time, 27 of them between midnight and 6am. If the logs are UTC rather than Central, that number shrinks a lot. This is the single biggest swing in the analysis, and it’s a compliance finding, so I’ve flagged it as an assumption everywhere it appears.

2. Does a call status of completed mean someone actually spoke to the customer? Why it matters: it’s the basis of the 43% connect rate. If it only means the call ended cleanly, then real contact is lower and the “Connected” bucket is overstated.

3. Which partners send by API and which by file transfer, and how often do the files land? Why it matters: 20 LeadBridge leads took 3–6 hours to first touch, all in one batch. If that’s a twice-daily file drop, it’s a delivery-schedule problem, not a dialer problem, and the fix is a partner conversation rather than engineering work.

4. What is the official outreach cadence, and who owns the system that schedules it? Why it matters: I inferred “call, then text about 15 minutes later, then a call about 6 hours later” from the dominant pattern. My 45 abandoned leads are defined against that inference. If the real policy is different, the count changes.

5. Does Legal treat a blank SMS consent field as consent? Why it matters: 17 leads have no consent recorded, and 6 of them were texted. I’ve treated blank as “no”. If that’s right, those 6 are breaches. If it’s wrong, I’ve overstated the exposure.

6. What is the conversation-to-close rate, and the average fee per close? Why it matters: I can show roughly 36 recoverable conversations, but I can’t turn that into dollars without these two numbers. They also decide whether this project outranks whatever else is competing for engineering time.

7. Are partners paid per lead, per closed deal, or on a flat fee? Why it matters: it sets the true cost of an untouched lead. If we pay per lead, the 18 untouched OfferNest leads cost real money on top of the lost revenue.

2. Topics to cover

The OfferNest batch. 18 consecutive lead IDs, 11–16 May, 60% of their volume in that window, while other OfferNest leads those same days were worked normally. I think one ingestion batch landed in the database but never reached the dialer queue. Does that match anything you saw? Did OfferNest raise it, or did we find it?

Abandoned cadences. 45 leads stopped after one or two touches, spread across 25 of 40 agents and all six partners, so it doesn’t look like individual behavior. My candidates are queue overflow, an expiring scheduling window, or lead reassignment. What’s actually most likely?

Duplicates. 10 people arrived through two partners each, 3 to 30 hours apart, worked by two different agents. Every pair has a *different property zip*. My read is these may be genuine second intents (selling in one city, buying in another), so we should merge the contact experience but keep both opportunities and both partners' attribution. Is that right, or should one be treated as a true duplicate and dropped? This decides whether a "fix" costs us a referral fee.

The speed assumption. The brief says odds drop sharply by the minute, and I believe it. But in this extract the 20 late-touched leads connected *more* often (65% vs 48%). Twenty leads proves nothing, and it may just reflect who those partners' customers are. Still: is speed-to-lead measured anywhere at Clever, or is it taken as received wisdom? If it isn't measured, I'd want to test it properly once the pipeline is reliable.

What I'd leave alone. The 42 leads that got the full cadence and never answered. Adding attempts there would raise compliance risk and muddy measurement. Also: shaving minutes off a first touch that's already inside SLA, replacing the dialer, and asking partners to standardize phone formats. I expect pushback on at least one of these.

The status field. Every one of the 206 leads still says new, including the 89 we spoke to. Is that a reporting artifact of the extract, or does the CRM genuinely never learn the outcome?

3. Questions about how Clever works

Things I'd ask in my first month regardless of this project:

- Who is accountable today when a lead goes untouched? If the answer is "nobody specifically", that's the real finding.
- How do Sales, Ops and Engineering currently agree on what happened? Is there a shared number at all?
- How do partners find out about problems: from us, or from their own customers?
- What have you already tried here, and what stopped it from sticking? I'd rather not re-run a failed experiment.
- Is there an existing dashboard or metric I should be extending instead of inventing something new?
- How much engineering time would realistically be available in month 1?
- Who has to agree before a metric becomes "the" number?

4. Questions about the role

- What does good look like for this role at 90 days?

- Which of these problems is already someone else's, and where would I be stepping on toes?
 - How does Product usually work with Ops here: do I bring the fix, or do we build it together?
 - What's the appetite for an engineering-facing PM? A good deal of my plan is instrumentation, which is unglamorous and easy to deprioritize.
-

If I only get 15 minutes

Cut to: the timezone question, the cadence-ownership question, the close rate and fee, the OfferNest batch, and the duplicate-identity decision. Those five change the most. The prototype can be a link they open later.

What I'd send beforehand

The proposal and the dashboard link, 24 hours ahead, with one line: "The seven questions in section 1 are where I'd most like your help; everything else can wait for the call."