

Every Lead Accounted For

Clever Real Estate: Senior Product Manager assessment


Joshua Agbabiaka · September 2026 · Fictionalized extract: 206 leads, 379 outreach attempts, 4-17 May 2026

The diagnosis. Clever can see successful component events - an API accepts a lead, a call is placed, a text is delivered - but not whether the lead received the promised experience. That is how every system can look green while customers are still missed. In this extract, only **56%** of leads received the expected handling, while **36%** were failed by the process rather than by customers who did not answer.

1. What the data says

Leads in, outreach out: where the two systems disagree

Every one of the 379 attempts maps to a known lead. The leakage is in treatment, not record linkage.



Outcome	Rule	Leads
Untouched	No outreach attempt at all	18
Dead-end	Attempts existed; no reachable channel succeeded	12
Abandoned	Started, never completed expected cadence	45
Exhausted	Full cadence completed; customer never responded	42
Connected†	At least one completed call	89

† Assumption: completed indicates a live conversation; confirm with Clever. **Separating Abandoned from Exhausted** prevents pipeline failure from being mistaken for poor lead quality.


Four findings drive the plan.

- 18 untouched leads sit in one OfferNest block. The leads exist, outreach does not; which handoff broke remains a hypothesis to validate against system logs.
- 45 abandoned leads span 25 of 40 agents and all six partners. Treat the issue as systemic; trace the scheduler before changing it.
- All 206 leads still read new, including 89 with completed calls. The business record never learns the outcome.
- 379 logged attempts produced 92 conversations. 32, all texts to landlines, never reached a human; 5 leads were called again after a completed call. Activity is not treatment.

The extract shows where the experience breaks, not every root cause. Discovery turns hypotheses into facts before large system changes.

2. What it's costing

What the leak is costing



3. Priorities: fix the experience, not the loudest symptom

First stop valuable leads from dying silently; then make failures diagnosable and recoverable; only then optimize.

Priority	Problem	Why	Action
P0	Consent + quiet-hour risk	Asymmetric downside; small guardrail	Ship guardrail; confirm rule with Legal
P0	Untouched OfferNest block	18 leads; partner-visible	Recover now; trace the handoff
P0 -> P1	Abandoned cadences	45 leads; largest recoverable pool	Instrument first; then fix root cause
P1	Landline dead-ends	12 leads; exact pattern	Add call fallback
P1	Duplicate / over-contact	10 people; 5 contacted after connect	Normalize identity; stop duplicate journey
Instrument	LeadBridge delay	Delay proven; cause unknown	Add timestamps before redesigning

Some duplicate people may represent legitimate second intents. Merge the contact journey, not the underlying opportunities or partner attribution.

What I would not do yet. I would not replace the dialer, telecom vendor, or CRM: the data proves outcome failure, not platform failure. I would not optimize first-touch minutes that already meet SLA, treat the 42 Exhausted leads as defects, or ask partners to standardize phone formats when normalization is our responsibility.

4. A shared metric system

Layer	Metric	Question it answers	Today
North star	Lead Handling Success Rate	Did Clever do everything it was supposed to do?	56%
Reliability	Process Failure Rate	How many leads did our process fail?	36%
Speed	First-Touch SLA	Did we act quickly enough?	76%
Follow-up	Cadence Completion	Did we finish the promised sequence?	48%
Business outcome	Connect Rate	Did we actually reach the customer?	43%
Guardrail	Critical Exceptions	Are we creating legal or customer risk?	65 leads*

Lead Handling Success = first touch inside SLA, and Clever either connected or completed the cadence. Connect Rate remains an outcome metric because it moves with lead quality. Guardrails prevent the north star from improving at the expense of compliance or customer experience.

* 65 unique leads carry at least one critical exception; flags can overlap.

5. The product: a Lead Lifecycle Ledger

The deliverable is not another dashboard. It is a trustworthy lead-state model that answers, for any lead: what should have happened, what did happen, what is late or broken, and what happens next.

The five stages, and what leaks at each

Dashed stages leave no record in Clever today, which is why these failures are invisible rather than merely unfixed.



Also at stage 4: 10 texts sent without SMS consent · up to 73 attempts outside 8am–9pm local, 27 between midnight and 6am* · 5 leads called after we had already spoken to them.

Fields to add. Lifecycle: ingested_at, routed_at, first_touch_at, outcome_at · Identity: normalized phone key, partner source, opportunity ID · Decision context: route destination, skip_reason, next expected action · Outcome: lead outcome, connection status, recovery status

Detection and recovery. Alert on absence: no touch within SLA, overdue follow-up, stopped partner feed. Reconcile every hop: sent -> ingested -> routed -> attempted. Maintain an exception queue with breach, age, assignment, and next action. Make recovery explicit: retry, call manually, correct identity, or close with an auditable reason.

The working prototype. A live ledger on this extract, backed by Postgres and visualized in Preset, with outcome logic in version-controlled SQL. It should show executive health, outcome distribution, concentration by partner/day, and an exception queue of leads currently stuck. The queue is what turns BI into an operational product.

Live dashboard: clever-lead-ledger.vercel.app · **How to read it:** github.com/Temitaio/clever-lead-ledger/blob/main/docs/dashboard-manual.md

All artifacts (diagrams, KPI taxonomy, SQL, dashboard code): github.com/Temitaio/clever-lead-ledger

6. Execution: from diagnosis to a durable operating system

From diagnosis to a durable operating system

Six phases. Timing stays visible so the sequence is testable, not aspirational.

■ I own / prototype ■ Engineering ■ Cross-functional

1. Discovery Week 1	Understand the real failure modes before prescribing changes	• Reconcile production data • Confirm SLA, cadence, consent rules	• Hand-trace a sample with Ops + Eng • Inspect ingestion to telecom path
2. Define Week 1-2	One shared definition of correct behaviour	• North star + lifecycle states • Alert thresholds, recovery rules	• Exception taxonomy • Ownership model
3. Delivery Weeks 2-4	Fix the top leaks; instrument the missing stages	• Lifecycle timestamps, skip reasons • Status writeback, identity normalization	• Landline to call fallback • Ledger v1 + exception queue
4. Launch Month 1	Get it into daily use, not just shipped	• Start with Ops + Product • Pilot partner scorecard internally	• Validate alerts against manual checks • Train Ops on exception handling
5. Monitor Month 1 on	Prove the system works, and that it is used	• Six metrics + time to detect / recover • Alert precision, mute rate	• Exception-queue action rate • Any surviving spreadsheet
6. Iterate Months 2-3	Automate only what proves repeatable	• Automate frequent recovery paths • Conversation-to-close attribution	• Partner delivery agreements • Clean speed-to-lead experiments

Effort is relative, not engineer-days: consent guard S | landline fallback S | lifecycle timestamps M | identity resolution M/L | cadence root cause unknown until traced | real-time ledger L

The week-1 deliverable is agreement, not a dashboard. Without one shared definition of a missed lead, every later fix becomes an argument about whose number is right. Two sequencing rules hold throughout: consent guardrails ship in week 1 regardless, because that exposure should not wait for a root cause; and the cadence scheduler is only changed after tracing confirms the mechanism.

Success criteria. No eligible lead breaches first-touch SLA without entering the exception queue · Abandoned and Dead-end become rare, diagnosable outcomes · every lead carries a truthful lifecycle status · vendor or platform changes happen only when instrumentation shows they are causal.

Open questions for the live session

- **Business rules:** What is the official first-touch SLA and cadence, and what does completed mean for a call?
- **Data semantics:** Are received_at and attempted_at in the same timezone? Is received_at partner-submit time or Clever-ingest time?
- **Architecture:** Which partners arrive by API vs. file, where does routing state live, and who owns the scheduler-to-dialer handoff?
- **Economics:** What are conversation-to-close rate and average referral value, so at-risk leads can be translated into dollars?
- **Compliance:** How should unknown SMS consent and state-specific quiet hours be treated?
- **Capacity:** If engineering time is constrained, which failure carries the highest customer and business cost?

Working principle. Clever does not appear to lack events; it lacks a shared lead-level outcome model that turns those events into an answer to the business question: did this lead get the right treatment, at the right time, and if not, why?